

E-Commerce Consumer Behavior Analysis

Business Analytics Consulting Report

BA Frameworks: PED · Vertical Analysis · Scenario Modeling · Market Structure

AI Tools: ThoughtSpot · WriteMyPRD

Dataset: 1,000 transactions · 28 attributes · 24 product categories

Executive Summary

This report analyzes 1,000 e-commerce transactions across 24 product categories to deliver consulting recommendations on pricing, promotion, market structure, and customer-loyalty strategy. Findings are structured against Business Analytics frameworks: Price Elasticity of Demand, Vertical Analysis, Scenario Modeling, and Market Structure classification.

Key findings:

- The top-selling category, Electronics, has a Price Elasticity of Demand of -1.366 — classified as Elastic. Demand is price-responsive, so margin protection (not price cuts) should be the lever.
- Vertical analysis reveals a long-tail revenue distribution: top 5 categories drive ~25% of revenue, with no single category exceeding 5.5%. Market is fragmented.
- Scenario modeling: a Marketing Push (+8% campaign cost) is projected to lift revenue by 1.43%, while a 15% Price Drop loses 3.38%. Recommend marketing investment over price war.
- Loyalty Discount rule (>\$500 spend OR Brand Loyalty=5) qualifies 208 customers (20.8%) costing \$5,315.77 — a manageable retention investment.
- Market is best classified as Monopolistic Competition: 24 categories, HHI of 429.5 (well below the 1500 competitive threshold), and meaningful brand-loyalty variance.

Q1 — Price Elasticity of Demand

Question

Calculate a proxy PED for the top-selling Purchase_Category by comparing Frequency_of_Purchase across Discount_Sensitivity segments.

Method


Top-selling category by units: Electronics (368 total units sold). Used Discount_Sensitivity segments as price-sensitivity proxies and applied the midpoint (arc) elasticity formula:

$$PED = (\% \text{ change in } Q) / (\% \text{ change in } P)$$

Where "Not Sensitive" customers represent the high-price segment and "Very Sensitive" the low-price (post-discount-equivalent) segment.

Results

Discount Sensitivity	Avg Price (\$)	Avg Quantity	N
Not Sensitive	215.405	7.524	21
Somewhat Sensitive	299.99	6.524	21
Very Sensitive	251.592	6.083	12



Verdict

PED = -1.366 → Elastic

Since $|PED| = 1.366 > 1$, demand for Electronics is price-responsive. A 10% price cut is projected to lift quantity by ~13.7%, but only if margins permit. Recommendation: avoid discounting; lead with bundling and value-add positioning.

Q2 — NumPy Demographic Analysis

Question

Calculate the average Time_Spent_on_Product_Research and standard deviation of Purchase_Amount across Gender, Income_Level, and Education_Level segments using NumPy.

Method

```
np.mean(arr) → average research time np.std(arr, ddof=1) → sample std dev of purchase amount
```

Results

By Gender

Gender	Avg Research Time (hrs)	Std Dev of Purchase (\$)	N
Agender	0.95	135.36	19
Bigender	0.89	133.44	20
Female	0.98	132.98	452
Genderfluid	1.12	137.24	17
Genderqueer	1.08	136.51	12
Male	1.05	129.6	449
Non-binary	0.94	120.53	16
Polygender	0.9	145.03	15

By Income_Level

Income_Level	Avg Research Time (hrs)	Std Dev of Purchase (\$)	N
High	1	131.1	515
Middle	1.03	132.12	485

By Education_Level

Education_Level	Avg Research Time (hrs)	Std Dev of Purchase (\$)	N
Bachelor's	1	126.54	341
High School	1.04	133.85	331
Master's	0.99	134.07	328

Insight

Standard deviation of purchase amount is remarkably consistent across all demographic cuts (\$120–\$145), suggesting purchase variability is driven more by category and intent than demographics. Average research time ranges 0.89–1.12 hours — a narrow band indicating most customers spend roughly an hour comparing options regardless of demographic. This argues for category-level personalization over demographic-level personalization.

Q3 — Vertical Analysis

Question


Show each Purchase_Category as a percentage of total revenue. Identify the "R&D-equivalent" category (smallest revenue share + high engagement).

Method

$$\text{Pct_of_Revenue}[\text{category}] = \text{Total_Revenue}[\text{category}] / \text{SUM}(\text{Total_Revenue}) \times 100$$

Top 10 Categories by Revenue Share

Category	Revenue (\$)	Transactions	% of Revenue
Jewelry & Accessories	\$15,139.36	50	5.50%
Sports & Outdoors	\$14,610.51	51	5.31%
Electronics	\$13,842.41	54	5.03%
Software & Apps	\$13,601.41	43	4.95%
Toys & Games	\$13,536.46	47	4.92%
Home Appliances	\$13,191.82	50	4.80%
Food & Beverages	\$12,966.96	44	4.71%
Packages)	\$12,731.16	43	4.63%
Health Care	\$12,149.03	41	4.42%
Gardening & Outdoors	\$11,782.28	44	4.28%



R&D-Equivalent Category

Arts & Crafts represents only 2.74% of total revenue — the "R&D-equivalent" — yet customers spend 1.09 hours researching it (above-average engagement). Like an R&D line item, it's a small revenue contributor today but absorbs disproportionate customer attention. This is either an emerging opportunity (worth more investment) or a poor product-market fit (worth pruning). Recommend running a focused customer-discovery sprint before deciding.

Q4 — GDP Impact on Demand Curve (Target: India)

Question

Research the GDP growth of a target market and explain how a 2% GDP shift would affect the demand curve for luxury vs utility categories in the dataset.

GDP Outlook — India FY2026

India is the fastest-growing major economy globally. Recent forecasts (April 2026 outlook):

- IMF: 6.6% (FY26), revised upward
- RBI: 7.3% (FY26), revised from 6.8%
- World Bank: 6.5% (FY26)
- Moody's: 6.4% (2026), fastest among G20

Consensus baseline: ~6.5–7.0% real GDP growth for FY26.

Demand Curve Theory

GDP shifts move the demand curve laterally — they do not change price elasticity within a category. A 2% upward GDP shift (to ~8.5–9% growth) injects disposable income into the economy; a 2% downward shift (to ~4.5–5%) compresses it. The magnitude of the curve shift depends on the income elasticity of demand (YED) for each category.

Impact on Dataset Categories

Category Type	Examples (from dataset)	YED (typical)	Demand Curve Response to +2% GDP
Luxury	Jewelry & Accessories, Electronics, Luxury Goods, Travel & Leisure	> 1 (income-elastic)	Shifts right strongly (+8–14% volume). Premium spend accelerates as discretionary income rises.
Utility	Groceries, Food & Beverages, Health Care, Baby Products, Animal Feed	< 1 (income-inelastic)	Shifts right modestly (+1–3% volume). Demand floor exists; consumption tied to need, not income.
Mid-tier	Books, Office Supplies, Mobile Accessories	~1 (unitary)	Shifts right proportionally (~2–4% volume).

Strategic Implication for Market Entry

In a high-GDP-growth scenario (+2%), prioritize inventory and marketing spend in Jewelry & Accessories, Electronics, and Travel — these will see the largest revenue lift. In a -2% downturn scenario, shift mix toward Groceries, Health Care, and Food & Beverages — these have demand floors and protect downside. Build merchandising flexibility to rotate between these poles based on quarterly GDP signals.

Q5 — Scenario Manager: Competitor Response

Question


Model two competitor-response scenarios on the dataset: (a) a 15% price drop, and (b) a marketing push that lifts engagement one level. Compare projected revenue.

Method

Scenario A applies a 15% reduction to Purchase_Amount and a Discount_Sensitivity-keyed quantity lift (Very Sensitive +25%, Somewhat +12%, Not Sensitive +3%). Scenario B applies an engagement-keyed quantity lift (None→+5%, Low→+10%, Medium→+15%, High→+8%) and subtracts an 8% campaign cost.

Results

Scenario	Projected Revenue	Δ vs Baseline	Δ %
Baseline (No Action)	\$275,063.88	+\$0	+0%
A: 15% Price Drop	\$265,754.01	-\$9,309.87	-3.38%
B: Marketing Push (+8% cost)	\$279,007.53	+\$3,943.65	+1.43%



Recommendation

RECOMMENDED: B: Marketing Push (+8% cost)

A reactive 15% price cut would destroy ~3.4% of revenue because the demand lift from price-sensitive customers does not offset the revenue lost on price-insensitive ones (a classic elastic-demand scenario where blanket cuts fail). The marketing push delivers a positive 1.4% lift after the campaign cost, and is sustainable rather than reactive — it builds brand equity rather than trading it away.

Q6 — Automated Dashboard: MAU vs Avg Spend

Question

Create an automated dashboard showing Monthly Active Users (proxy for MAU, derived from Time_of_Purchase) vs Average Purchase Amount per Customer (proxy for CAC efficiency) using sparklines and conditional formatting.

Deliverable

Built in Ecommerce_BA_Dashboard.xlsx (Dashboard sheet). Features:

- KPI cards (Total Revenue, Customers, Avg Ticket, Loyalty Discount Issued)
- Monthly metrics table with sparklines for MAU and Avg Spend trends (Excel 2010+ native sparklines, embedded via XML)
- Conditional formatting (3-color scale red→yellow→green) on MAU and Avg Spend columns
- Live Excel chart of MAU trend, regenerable from underlying data
- SUM/AVERAGE formulas at the totals row — fully dynamic

Monthly Snapshot

Month	MAU	Total Revenue (\$)	Avg Spend (\$)
2024-01	75	19,089.22	254.52
2024-02	68	20,294.28	298.45
2024-03	93	27,409.79	294.73
2024-04	100	28,617.22	286.17
2024-05	76	18,658.99	245.51
2024-06	89	25,253.23	283.74
2024-07	94	25,260.62	268.73
2024-08	98	27,473.39	280.34
2024-09	84	24,054.76	286.37
2024-10	79	19,492.38	246.74
2024-11	80	21,602.54	270.03
2024-12	64	17,857.46	279.02

Q7 — 6-Month Forecast

Question

Use a forecasting method to project monthly purchase volume for the next 6 months based on historical Time_of_Purchase data.



Method

Holt damped-trend exponential smoothing (statsmodels)

Statsmodels' Holt-Winters with damped additive trend was selected over standard linear regression because it adapts to the level + trend structure of the 12-month history without overfitting seasonal patterns from a single year of data. Forecasts are floored at 85% of historical minimum to prevent unrealistic negative drift.

6-Month Forecast

Month	Forecast Transactions
2025-01	85
2025-02	85
2025-03	85
2025-04	85
2025-05	85
2025-06	85



Interpretation

The damped-trend forecast settles around ~85 transactions/month, very close to the historical mean of 83. With only 12 months of history there is insufficient signal to detect annual seasonality reliably. Recommendation: re-run the forecast with `seasonal_periods=12` once 24+ months of data are available; until then, treat $\pm 15\%$ as the prudent confidence band.

Q8 — Loyalty Discount Rule

Question

Write an if-else statement assigning a 10% Loyalty Discount to customers spending >\$500 OR with Brand_Loyalty = 5.

Implementation

```
def loyalty_discount(row):    if row["Purchase_Amount"] > 500 or  
row["Brand_Loyalty"] == 5:    return round(row["Purchase_Amount"] * 0.10, 2)  
else:                        return 0.0  
df["Loyalty_Discount"] = df.apply(loyalty_discount,  
axis=1)
```

Equivalent Excel formula (used in the workbook):

```
=IF(OR(Purchase_Amount > 500, Brand_Loyalty = 5), Purchase_Amount * 10%, 0)
```

Impact

Metric	Value
Qualifying Customers	208 of 1000
% of Customer Base	20.8%
Total Discount Issued	\$5,315.77
Avg Discount per Qualifier	\$25.56

At 20.8% qualifying base and ~\$25.50 average discount per qualifier, total program cost equals roughly 1.9% of revenue — well within typical retention budgets (industry benchmark: 3–7%). The rule appropriately rewards high-value spenders AND brand-loyal repeat customers, rather than only one signal.

Q9 — ThoughtSpot Natural-Language Query

Question

Use ThoughtSpot to answer: "Which Income_Level and Customer_Loyalty_Program_Member combination has the highest average Customer_Satisfaction and lowest Return_Rate?"

Tool: ThoughtSpot Search Sage

ThoughtSpot Search Sage parses natural-language questions against an indexed dataset and returns auto-generated visualizations. The query workflow is:

- Upload Ecommerce_Consumer_Behavior_Analysis_Data.csv as a worksheet
- Tag Income_Level and Customer_Loyalty_Program_Member as ATTRIBUTES
- Tag Customer_Satisfaction and Return_Rate as MEASURES (aggregation: average)
- Type the natural-language query into the Search bar

Natural-Language Query

average customer_satisfaction and average return_rate by income_level and customer_loyalty_program_member sorted by customer_satisfaction descending

Equivalent SQL (auto-generated)

```
SELECT Income_Level, Customer_Loyalty_Program_Member,
AVG(Customer_Satisfaction) AS avg_satisfaction, AVG(Return_Rate) AS
avg_return_rate FROM ecommerce_consumer_behavior GROUP BY Income_Level,
Customer_Loyalty_Program_Member ORDER BY avg_satisfaction DESC;
```

Result (computed from the dataset)

Runs of the equivalent SQL on the dataset return the following segmentation. The most retention-friendly segment (highest satisfaction + lowest return rate) is then identified as the recommended target for retention campaigns.

Income Level	Loyalty Program	Avg Satisfaction	Avg Return Rate	N
Middle	Non-member	5.66	0.956	249
Middle	Member	5.54	0.983	236
High	Member	5.44	0.941	255
High	Non-member	4.98	0.938	260

Winning segment: Middle income, Non-Member (Avg Satisfaction = 5.66, Return Rate = 0.956). This is the most retention-friendly segment — invest in expanding it via lookalike acquisition and loyalty-program enrollment campaigns.

Q10 — WriteMyPRD: Dynamic Discount Engine BRD

Question

Use WriteMyPRD to generate a Business Requirement Document for a Dynamic Discount Engine that adjusts discounts in real time based on Discount_Sensitivity, Brand_Loyalty, and Purchase_Intent.

Tool: WriteMyPRD

WriteMyPRD is a GPT-powered PRD/BRD generator. Workflow: provide product name, problem statement, target user, key features, and success metrics — the tool produces a structured BRD. Below is the BRD generated using this template.

Business Requirement Document

Product Name: Dynamic Discount Engine (DDE)

Owner: E-Commerce Pricing Team

Date: April 2026

1. Problem Statement

Current pricing applies blanket discounts that erode margin on price-insensitive customers (who would pay full price) while under-incentivizing price-sensitive customers (who could be converted with sharper, targeted offers). Analysis shows $|PED| = 1.37$ in our top category, indicating elastic demand — yet our discounts are flat. We are leaving margin on the table on one side and conversion on the other.

2. Objectives

- Increase gross margin by 200–400 bps without reducing total transaction count
- Personalize discounts using three signals already captured: Discount_Sensitivity, Brand_Loyalty, Purchase_Intent
- Operate in real time at checkout (latency < 200ms)
- A/B-testable: every offer logged with treatment vs control flag

3. Target Users

- PRIMARY: All checkout users (1,000+ daily) — the engine fires automatically
- SECONDARY: Pricing analysts who tune the rule weights via an admin UI
- TERTIARY: Finance team who monitors margin impact via dashboard

4. Functional Requirements

FR-1: Real-time scoring service

On each checkout request, compute a Discount Score:

$$\text{discount_score} = w1 * \text{sensitivity_factor} + w2 * \text{loyalty_factor} + w3 * \text{intent_factor}$$

Where:

- sensitivity_factor: 0.0 (Not Sensitive) → 0.5 (Somewhat) → 1.0 (Very Sensitive)
- loyalty_factor: $1 - (\text{Brand_Loyalty} / 5)$ — inverse of loyalty (lower loyalty = bigger discount needed to convert)
- intent_factor: 0.2 (Need-based) → 0.5 (Wants-based) → 1.0 (Impulsive)
- Default weights: $w1=0.5, w2=0.3, w3=0.2$ — tunable by Pricing team

FR-2: Discount mapping table

Map score → discount %:

Score Range	Discount Applied
Score 0.0–0.3	0% (no discount, full margin)
Score 0.3–0.5	5%
Score 0.5–0.7	10%
Score 0.7–1.0	15% (max)

FR-3: Override layer

- Loyalty Program Members get a +2% floor regardless of score (rewarding membership)
- Brand_Loyalty = 5 customers capped at 5% discount (don't over-discount your superfans)

FR-4: Admin UI

- Tune w1/w2/w3 weights in real time
- Adjust score-to-discount mapping bands
- Toggle overrides on/off

FR-5: Logging & analytics

- Every offer logged: customer_id, score components, applied %, accepted/declined, timestamp
- Daily margin impact report vs control (10% holdout)

5. Non-Functional Requirements

- NFR-1: Latency < 200ms p95 at checkout
- NFR-2: 99.9% uptime; if engine is down, fall back to flat 5% legacy discount
- NFR-3: GDPR/PII compliant — no raw demographic fields stored in scoring logs
- NFR-4: Audit trail of all weight changes

6. Success Metrics

- NORTH STAR: Gross margin % vs flat-discount baseline (target: +200 bps in 90 days)
- Conversion rate at checkout (must not decline >2%)
- AOV (Average Order Value) — target +5%
- Customer satisfaction (post-purchase NPS) — must stay within 1 point of baseline

7. Out of Scope

- Dynamic pricing of base product price (this is discount-only)
- Multi-currency / international rollout (Phase 2)
- Subscription pricing (different surface)

8. Dependencies

- Real-time customer feature store (Discount_Sensitivity must be precomputed and accessible <50ms)
- Checkout API surface — needs new discount_engine endpoint
- Finance approval on margin-erosion guardrails

9. Risks

- R1: Discounts erode brand premium — MITIGATION: cap loyal customers at 5%

- R2: Customers reverse-engineer the rule and game it — MITIGATION: add randomized noise ($\pm 2\%$) to discount, refresh weights monthly
- R3: Latency budget exceeded under peak load — MITIGATION: precompute scores nightly, refresh on session start

10. Timeline

- Weeks 1–2: Feature store integration + scoring service prototype
- Weeks 3–4: Admin UI + checkout API integration
- Weeks 5–6: A/B test setup + 10% rollout
- Weeks 7–10: Full rollout, monitoring, weight tuning

Q11 — Market Structure Classification

Question

Define the market structure: Oligopoly, Monopolistic Competition, or Perfect Competition?

Method — Three Tests

Three structural metrics determine market type:

- Number of competitors (categories as proxy for product variants)
- Concentration: Herfindahl-Hirschman Index (HHI) on revenue share
- Differentiation: variance in Brand_Loyalty and Product_Rating

Results

Test	Value	Interpretation
Number of Categories	24	Many (>10) → low concentration
HHI (revenue concentration)	429.5	<1500 = competitive, 1500–2500 = moderate, >2500 = concentrated
Brand Loyalty Mean / Std	3.03 / 1.42	Mid-range mean + high std → meaningful differentiation
Product Rating Mean / Std	3.03 / 1.44	Quality varies meaningfully across products

Classification

MONOPOLISTIC COMPETITION

Many categories (24) with low revenue concentration ($HHI < 1500$) and meaningful brand-loyalty variation indicates many sellers offering differentiated products — the textbook definition of monopolistic competition.

What this means strategically:

- Pricing power exists but is bounded — small premiums sustainable, big premiums penalized
- Branding and differentiation are how you win, not cost leadership
- Entry barriers are low — new competitors will appear; retention matters more than acquisition
- Product variety is the moat — investing in long-tail SKU coverage pays off

Q12 — Seaborn Boxplot: Transaction Volume by Payment Method

Question


Create a Seaborn boxplot showing the distribution of Purchase_Amount by Payment_Method, segmented by Device_Used_for_Shopping.

Code

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.boxplot(data=df,
            x="Payment_Method", y="Purchase_Amount",
            hue="Device_Used_for_Shopping", palette="Set2")
plt.title("Purchase Amount Distribution by Payment Method & Device")
plt.show()
```

Visualization



Insights

- Distributions overlap heavily — payment method is NOT a strong predictor of transaction size in this dataset
- Medians cluster around \$250–\$280 across all combinations
- Smartphone purchases show slightly tighter IQRs across payment methods, suggesting more disciplined mobile spending behavior
- No payment method has a meaningfully different upper tail — high-value transactions occur across all five payment types
- Operational implication: do not segment fraud-detection thresholds by payment method alone — the distributional signal is too weak

Conclusion

Twelve questions, three deliverable types, two AI tools, and one consistent narrative: this is a fragmented, monopolistically-competitive market where price elasticity is high, differentiation matters more than discount depth, and retention beats acquisition. The recommended path forward is a Marketing Push (Q5) over a Price War, a Dynamic Discount Engine (Q10) to personalize value capture, and continued investment in the high-engagement / low-revenue categories identified in the vertical analysis (Q3) before pruning them.

All numerical analysis is reproduced live in the accompanying Excel workbook (Ecommerce_BA_Dashboard.xlsx). Charts are sourced from the analysis script and embedded throughout this report.